



Big Data - Map Reduce

Dr. Matt Zhang
ITU



Outline

- ❑ Motivation
- ❑ How M/R Works
 - ❑ Programming Model
- ❑ Jobs
- ❑ M/R Types and Formats
- ❑ M/R Features
- ❑ Develop M/R Application
 - ❑ Word Count Example



Motivation

- Data-intensive computing has arrived:
 - Both user-facing services and batch data processing
 - Data analysis is key
- Need massive scalability and easy parallelism
 - PB's of data, millions of files, 1000's of nodes, millions of users.

Motivation (cont.)

- Need to do this cost effectively and reliably
 - Data warehouse too expensive
 - Teradata maintenance support fee > millions of \$ per year
 - Use commodity hardware where failure is the norm
 - Share resources among multiple projects

MapReduce to the rescue!

MR Features

- Simple data-parallel programming model and framework:
 - Designed for scalability and fault-tolerance
 - Automatic parallelization and distribution
 - Status and monitoring tools
 - Abstracts all the internal work away from developers
 - Can focus simply on writing Map and Reduce functions

MR Applications

- Pioneered by Google

- ☐ Processes 20 PB of data per day
- ☐ Popularized by open-source Hadoop project
- ☐ Used at Yahoo!, Facebook, Amazon, etc.

MR Applications

- **Google:** index construction for search, article clustering for Google news, statistical machine translation
- **Yahoo!:** web-map and spam detection for Yahoo! mail
- **Facebook:** Ad optimization and spam detection

MR and Data Analytics

□ Data cleaning:

- ❖ Preprocess data in order to reduce noise and handle missing values – goal is to improve learning

□ Relevance analysis:

- ❖ Remove the irrelevant or redundant attributes using correlation analysis

MR and Data Analytics

- Data transformation and reduction:
 - ❖ Generalize to higher-level concepts, and/or
 - ❖ Normalize data (an attribute value is scaled to be between 0.0 – 1.0), especially if neural networks or distance measurements are used in the learning step.
 - ❖ Data can be Reduced by applying methods ranging from wavelet transformation to discretization techniques

Programming Model

- APIs: <http://hadoop.apache.org/docs/r2.2.0/api/org/apache/hadoop/mapred/>

- **Input data type:** file of K/V records
- **Map function:** $(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$
- **Reduce function:** $(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$
- **Example:**


```
def mapper(line):
```

```
    foreach word in line.split();
```

```
        output(word, 1);           // (Kinter, Vinter)
```

```
def reducer(key, values):
```

```
    output(key, sum(values));      // list(Kout, Vout)
```



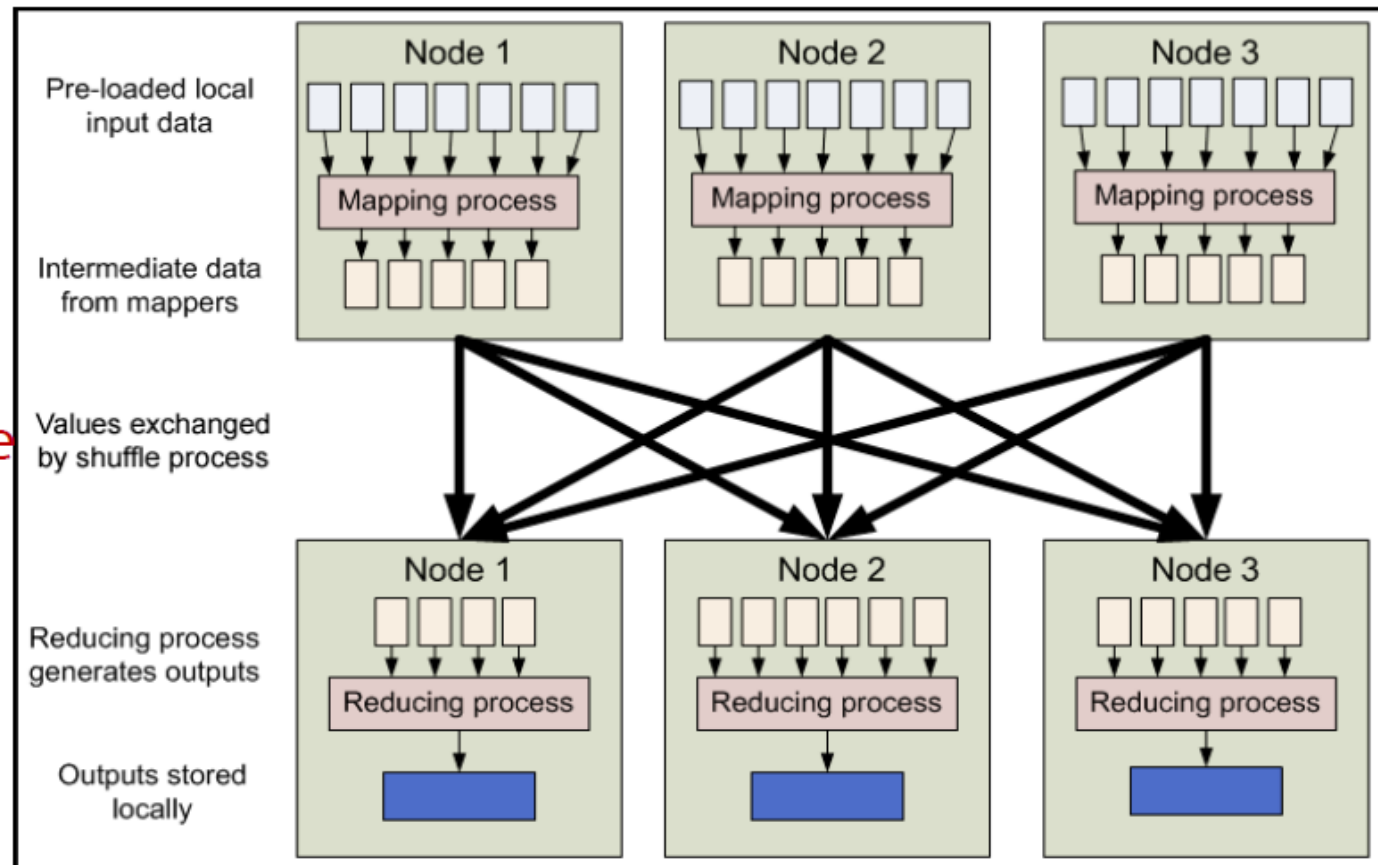
Programming Model

- **Push:** input split into large chunks and placed on local disks of cluster nodes
- **Map:** chunks (map tasks) are served to “mapper”
 - Prefer mapper that has data locally
 - Mappers save outputs to local disk before serving them to reducers; allows recovery
- **Reduce:** reducers execute reduce tasks only when map phase is completed

Programming Model

The Big Picture

Disk I/O
Intensive



High degree of parallelism: Master/Slave architecture

Partitioning

- **Partitioning/Shuffling**: divide intermediate key space across reducers:
 - K reduce tasks → K partitions (simple hash function)
 - E.g., K = 3, keys {1,2}, {3,4}, {5,6}

JobTracker

- Software daemons control MR jobs
- Resides on master node
 - Client submit MR jobs to JobTracker
 - It assigns MR tasks to other nodes
- Each slave node has a TaskTracker daemon
 - Instantiating the M/R tasks
 - Status reporting to JobTracker

Job Terminology

- Job

- user program

- Task

- Execution of a single M/R over a slice of data
 - If one fails, JobTracker will start another one
 - Speculative execution

Mapper

- Reads data of key/value pairs
- Run in parallel, each processing a portion of input
- Output also key/value pairs
- Mappers run on nodes with data locality
 - Minimize network traffic

Reducer

- Process the output intermediate key/value pairs from Mapper, and output results
- Intermediate values for each key are combined into a list
 - Same key goes to same reducer
 - Sorted key order – “shuffle and sort”
- Output zero or more final key/value pairs
 - Write to HDFS

Shuffling

- Shuffle and sort:

- All mappers typically have all intermediate keys

- All to all communication

- Bottleneck?

- Reducers cannot start until all Mappers finish

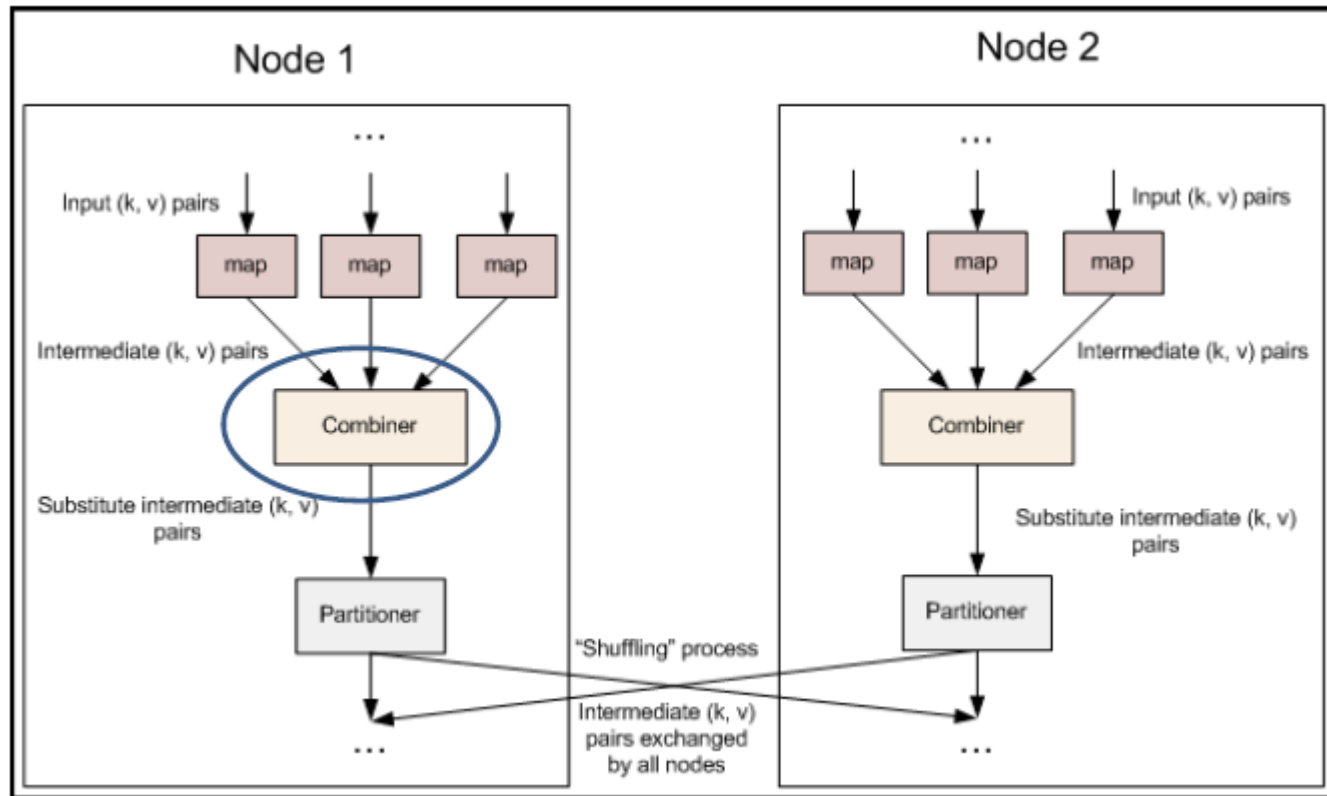
- In practice, Mappers transfer data after finishing

- avoids data transfer at the same time

Combiner

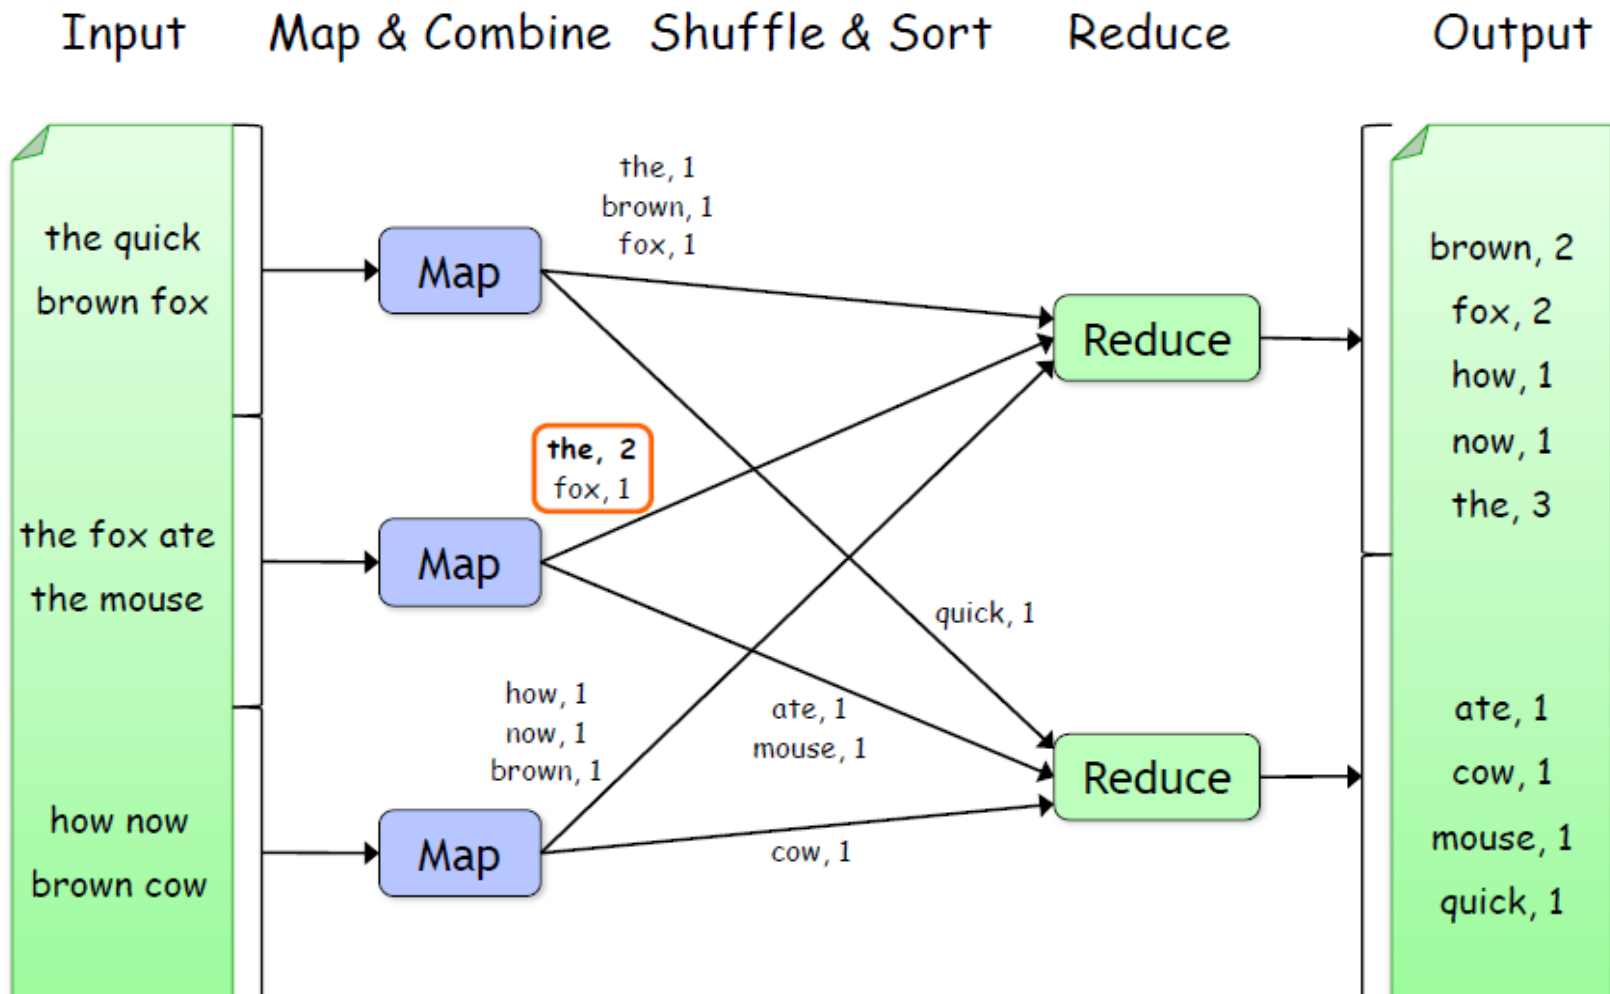
- Mappers produce large amount of intermediate data
 - Lots of network traffic
- Pre-aggregation
 - Mini-reducer
 - Runs on single Mapper's output
 - Input/output data types for Combiner/Reducer must be the same
 - Code is often the same as reducer

Combiner



- A combiner is a local aggregation function for repeated keys produced by the same mapper

Combiner Example



Word Count with Combiner

Fault Tolerance in MapReduce

❑ If a task crashes:

- ❖ Retry on another node
- OK for a map because it has no dependencies
- OK for reduce because map outputs are on disk

❑ If a node crashes:

- ❖ Re-launch its current tasks on other nodes
- ❖ Re-run any maps the node previously ran to get output data

Fault Tolerance in MapReduce (cont.)

- ❑ If a task is going slowly:
 - ❖ Launch second copy of task on another node (“speculative execution”)

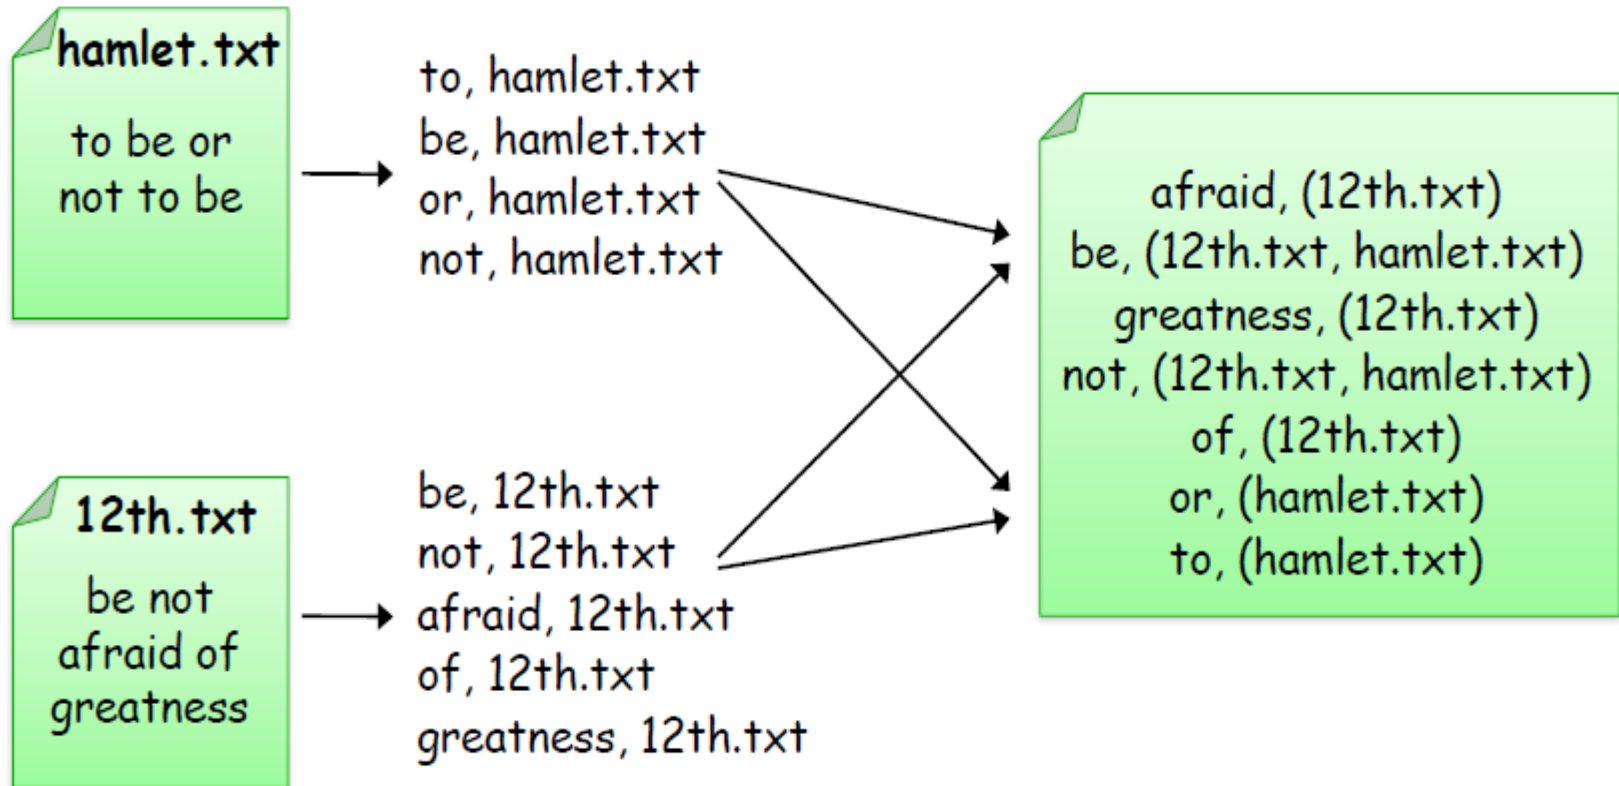
Example: Inverted Index

- ❑ **Input:** (filename, text) records
- ❑ **Output:** list of files containing each word

- ❑ **Map:**
 foreach word in text.split():
 output(word, filename)

- ❑ **Reduce:**
 def reduce(word, filenames):
 output(word, sort(filenames))

Example: Inverted Index



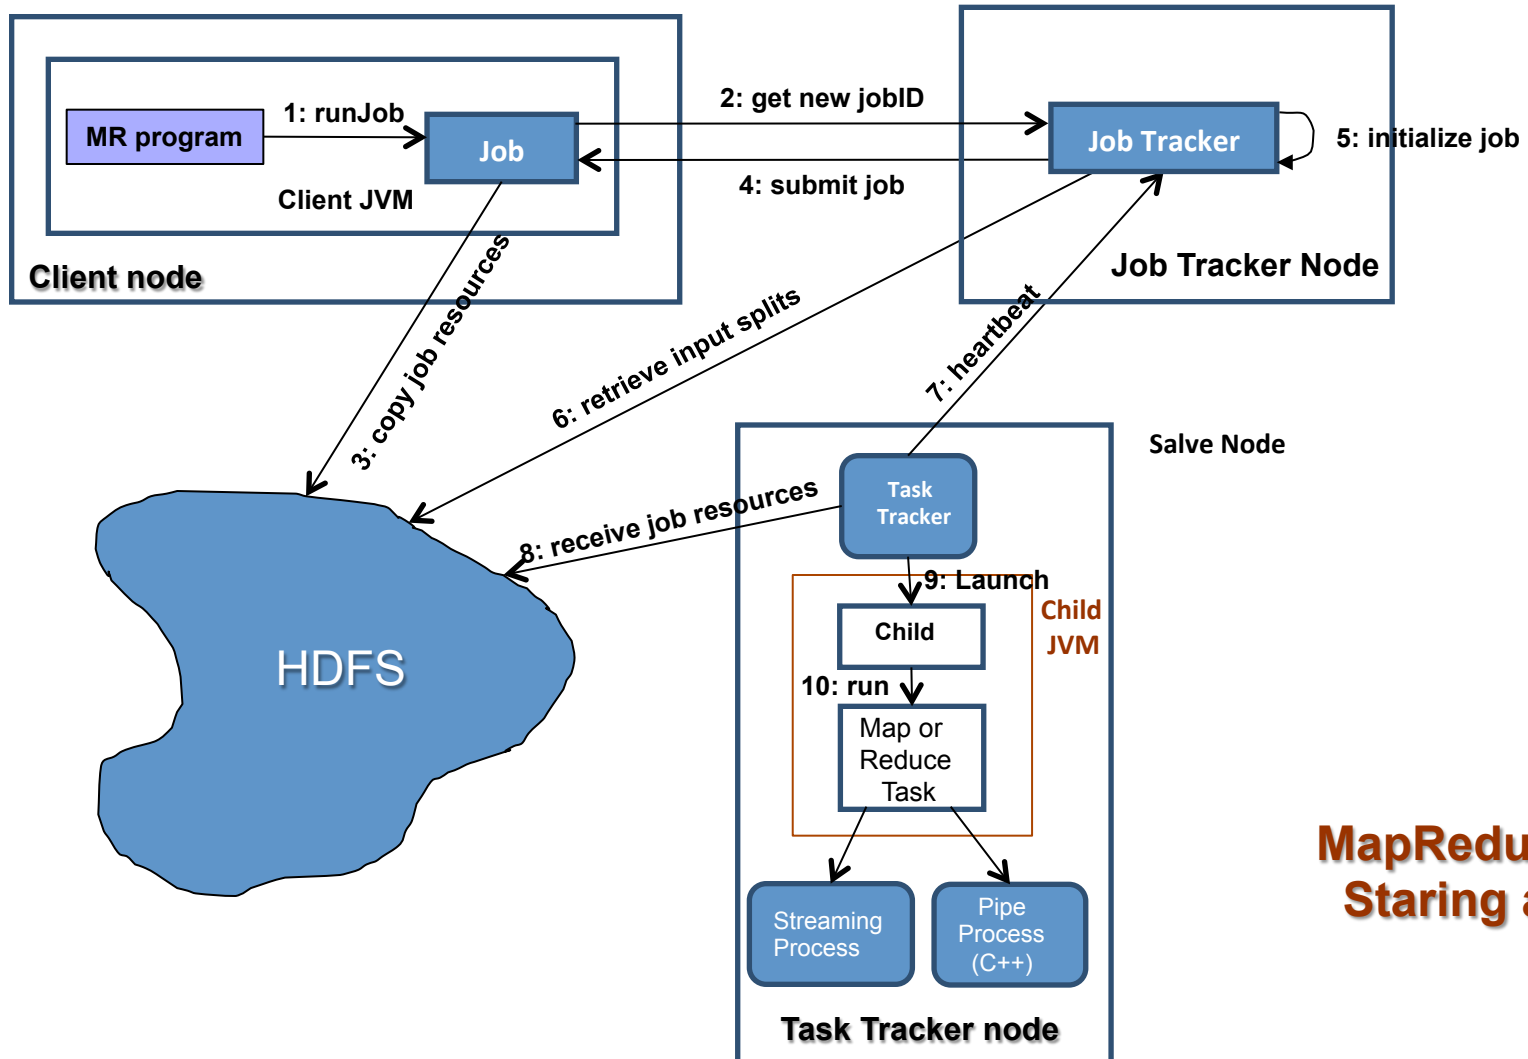
Inverted Index Example



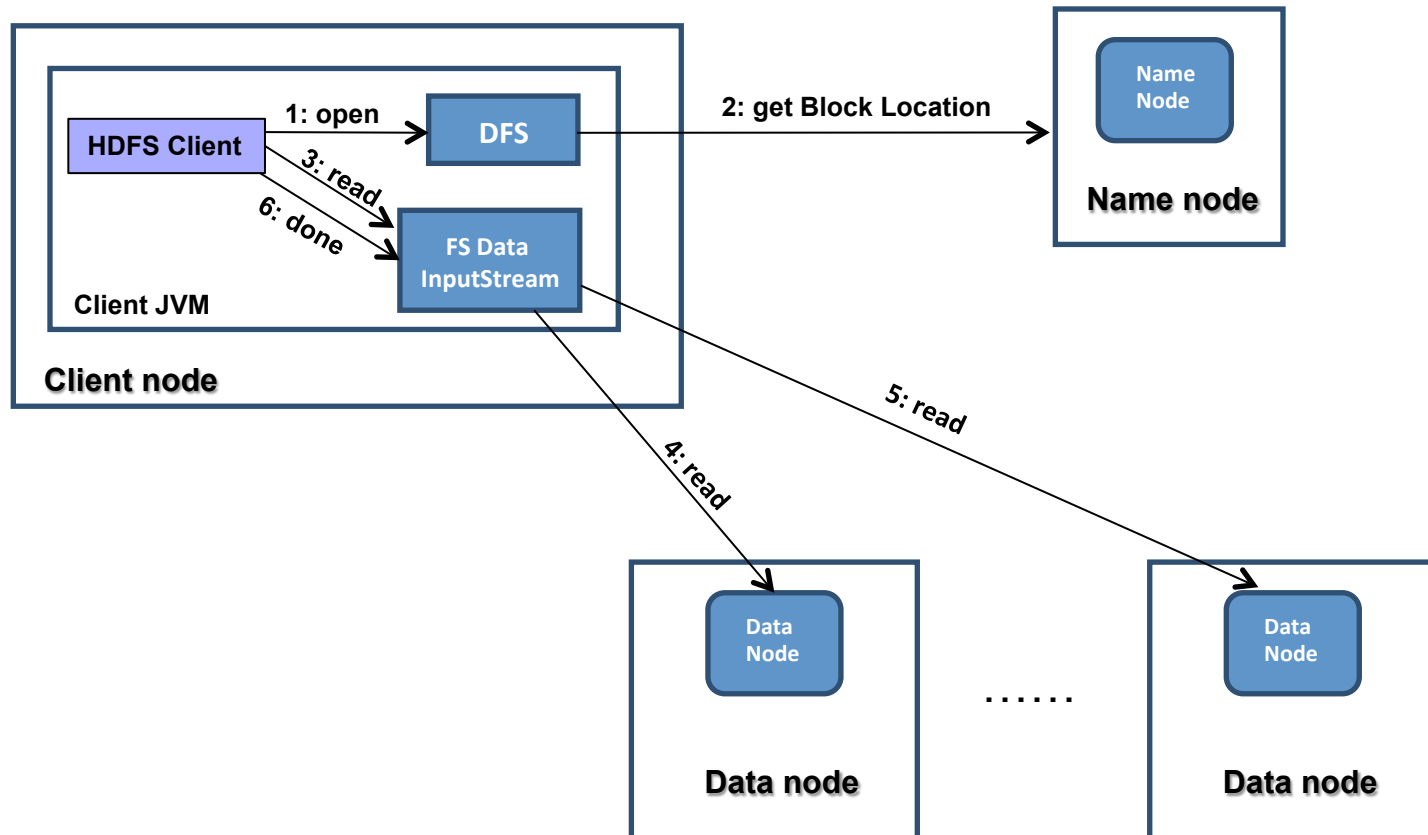
Applications

- What applications may perform well?
 - ❑ Modest computing relative to data
 - ❑ Data-independent processing of maps
 - ❑ Data-independent processing of keys
 - ❑ Smaller ballooning of map output relative to input

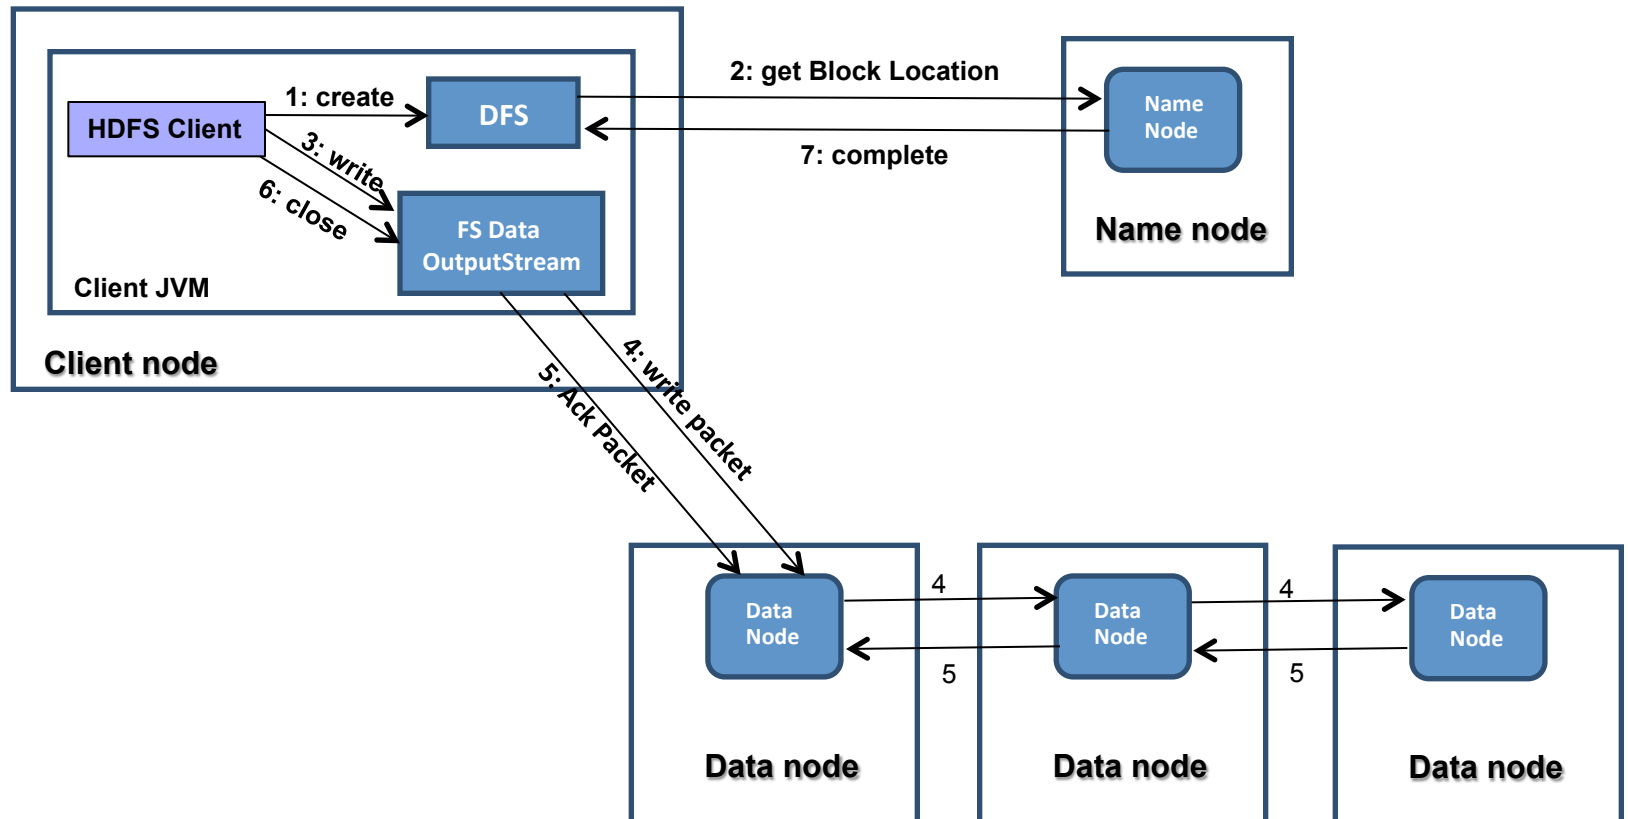
Starting a Job



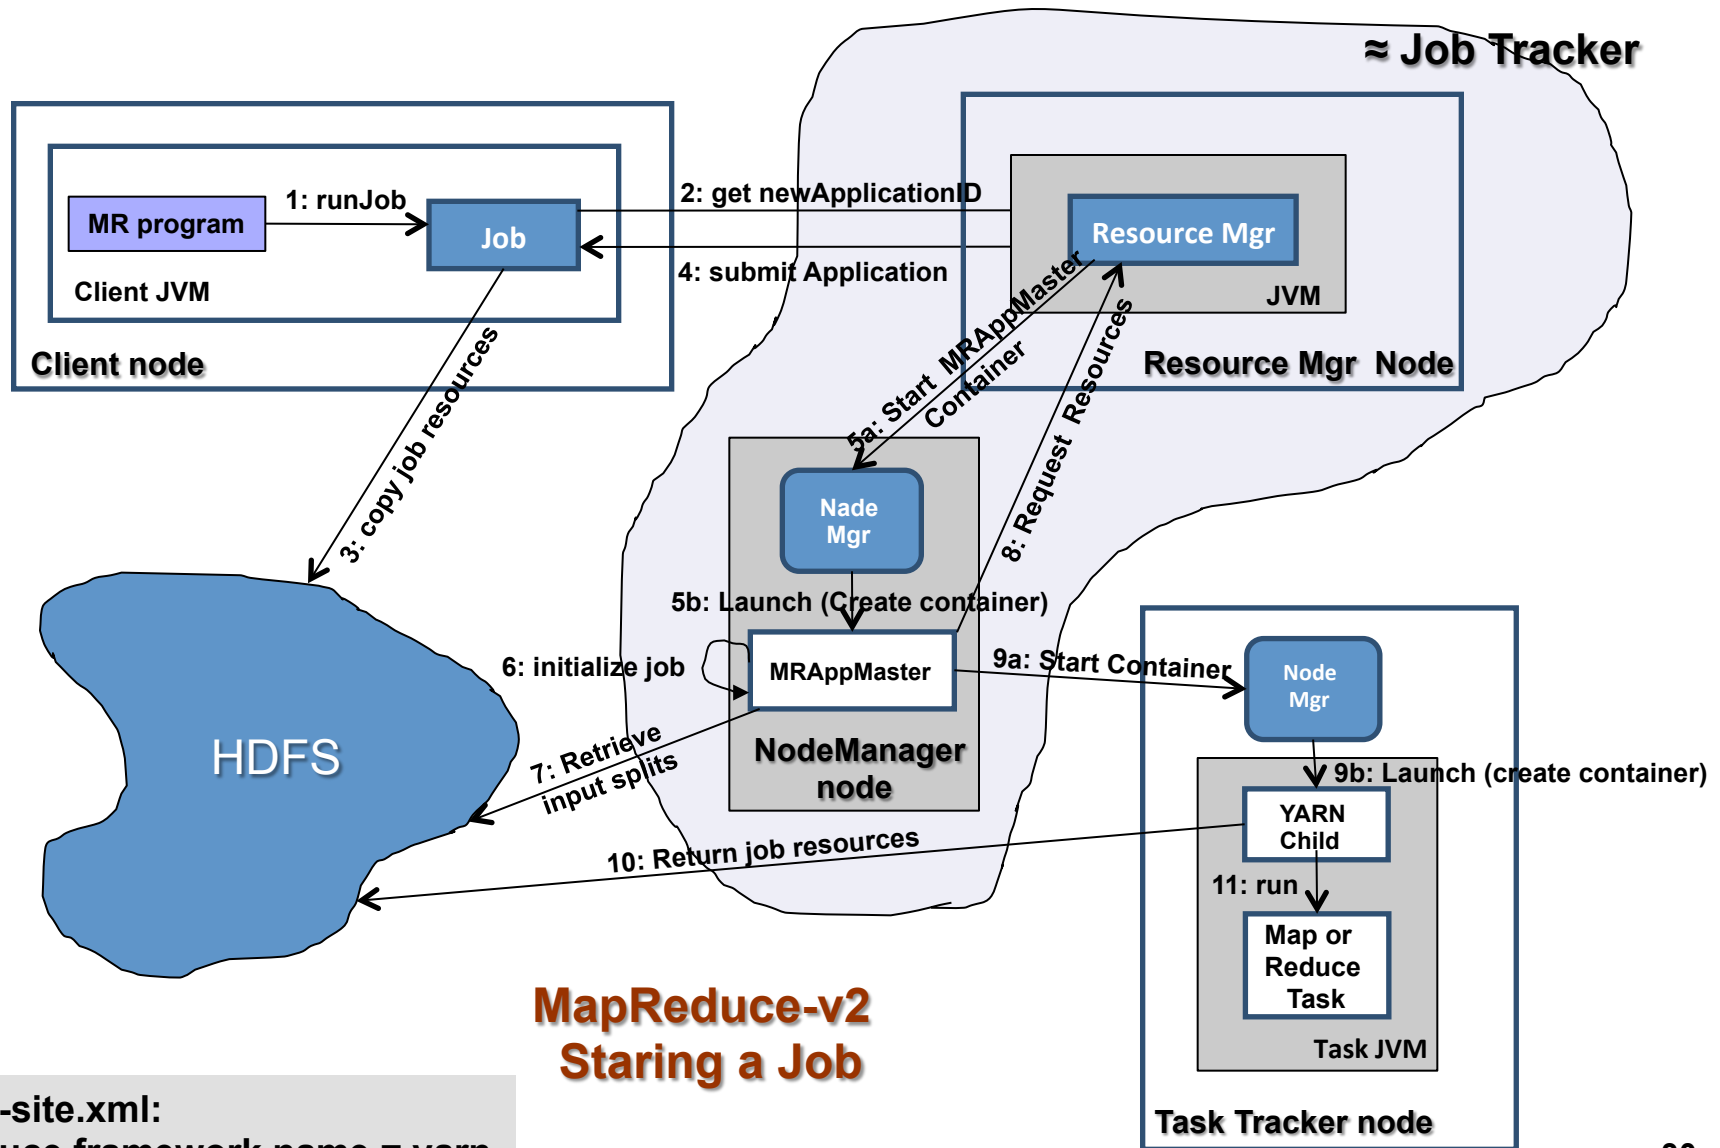
Read Anatomy



Write Anatomy



Starting a Job (Yarn)



Mapred-site.xml:
Mapreduce.framework.name = yarn



Developing M/R Applications

Developing MapReduce Application

- Write Map and Reduce functions and test them independently. **MRUnit** (<http://incubator.apache.org/mrunit>) is a library used to test the mapper() or reducer() as stand-alone function.
- **MRUnit** is used with **JUnit** to test MR Jobs as part of your **IDE** environment.
- Write a driver program to run a job.

Developing MapReduce Application

- Run the job from your IDE using a small subset of the data
- Debug using the IDE debugger.
- Run against the full dataset and in a cluster environment
- May expose issues that did not show up in the IDE testing.

Developing MapReduce Application

- After the program is working in a cluster, it is time for tuning through profiling.
- Before developing MapReduce job, we need to set up and configure the development environment.
- For details refer to Chapter-10, Hadoop: The Definitive Guide, 4th Edition

Word Count Example: Mapper

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class WordCountMapper extends MapReduceBase
    implements Mapper <LongWritable, Text, Text,
    IntWritable>
{
    // hadoop supported data types
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
```

Word Count Example: Mapper

```
// map method that performs the tokenizer job and  
// framing the initial key value pairs
```

```
public void map(LongWritable key, Text value,  
OutputCollector<Text, IntWritable> output, Reporter  
reporter) throws IOException
```

```
{
```

```
    // taking one line at a time and tokenizing the same  
    String line = value.toString();  
    StringTokenizer tokenizer = new  
        StringTokenizer(line);
```

Word Count Example: Mapper

```
// iterating through all the words available in that line  
// and forming the key value pair
```

```
while (tokenizer.hasMoreTokens())
```

```
{
```

```
    word.set(tokenizer.nextToken());
```

```
    // send to output collector which in turn passes the
```

```
    // same to reducer
```

```
    output.collect(word, one);
```

```
}
```

```
}
```

```
}
```

Word Count Example: Reducer

```
import java.io.IOException;  
import java.util.Iterator;  
  
import org.apache.hadoop.io.*;  
import org.apache.hadoop.mapred.*;
```

Word Count Example: Reducer

```
public class WordCountReducer extends  
    MapReduceBase implements Reducer<Text,  
    IntWritable, Text, IntWritable>  
{  
    // reduce method accepts the Key Value  
    // pairs from mappers, do the aggregation  
    // based on keys and produce the final output
```

Word Count Example: Reducer

```
public void reduce(Text key,  
Iterator<IntWritable> values,  
OutputCollector<Text, IntWritable> output,  
Reporter reporter) throws IOException
```

```
{
```

```
    int sum = 0;
```

```
    /* iterates through all the values available  
       with a key and add them together and  
       give the final result as the key and sum  
       of its values */
```


Word Count Example: Reducer

```
while (values.hasNext())
```

```
{
```

```
    sum += values.next().get();
```

```
}
```

```
output.collect(key, new IntWritable(sum));
```

```
}
```

```
}
```

Word Count Example: Driver

```
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.conf.*;  
import org.apache.hadoop.io.*;  
import org.apache.hadoop.mapred.*;  
import org.apache.hadoop.util.*;
```

Word Count Example: Driver

```
public class WordCount extends  
Configured implements Tool {  
    public int run(String[] args) throws  
Exception  
    {  
        //creating a JobConf object and  
        // assigning a job name for  
        // identification purposes  
        // Class JobConf -
```

Word Count Example: Driver

[/*http://hadoop.apache.org/docs/r2.3.0/api/
org/apache/hadoop/mapred/JobConf.html */](http://hadoop.apache.org/docs/r2.3.0/api/org/apache/hadoop/mapred/JobConf.html)

```
JobConf conf = new  
    JobConf(getConf(),  
        WordCount.class);
```

```
conf.setJobName("WordCount");
```

Word Count Example: Driver

```
// Setting configuration object with the  
// Data Type of output Key and Value  
conf.setOutputKeyClass(Text.class);  
conf.setOutputValueClass(  
    IntWritable.class);
```

Word Count Example: Driver

// Providing the mapper and reducer

// class names

```
conf.setMapperClass(
```

```
    WordCountMapper.class);
```

```
conf.setReducerClass(
```

```
    WordCountReducer.class);
```

Word Count Example: Driver

```
// the hdfs input and output directory  
// to be fetched from the command  
// line.
```

// Class Path - <https://hadoop.apache.org/docs/r2.2.0/api/org/apache/hadoop/fs/Path.html>

```
FileInputFormat.addInputPath(  
    conf, new Path(args[0]));  
FileOutputFormat.setOutputPath(  
    conf, new Path(args[1]));
```

Word Count Example: Driver

```
JobClient.runJob(conf);
```

```
/* Class JobClient –
```

```
    http://hadoop.apache.org/docs/r2.2.0/  
    api/org/apache/hadoop/  
    mapred/JobClient.html */
```

```
return 0;
```

```
}
```


Word Count Example: Driver

```
public static void main(String[] args)
    throws Exception
    {
        int res = ToolRunner.run(
            new Configuration(),
            new WordCount(), args);
        System.exit(res);
    }
}
```



Summary

- We introduced the programming model of MapReduce
- We discussed the components of MapRudece
- We covered the steps recommended for developing MapReduce Applications
 - Word Count Example