



Big Data - Hadoop Admin

Dr. Qing “Matt” Zhang
ITU

Checking HDFS Status

- *hdfs fsck* checks for missing or corrupt data blocks
 - Unlike Linux system fsck, it does not attempt to repair errors
- Can list all files, all blocks for each file, all block locations, or all racks
- Examples:
 - *hdfs fsck /*
 - *hdfs fsck / -files*
 - *hdfs fsck / -files -blocks*
 - *hdfs fsck / -files -blocks -locations*
 - *hdfs fsck / -files -blocks -locations -racks*

Checking HDFS Status

- Good idea to run *hdfs fsck* as a regular cron job that emails the results to administrators
- Choose a low-usage time to run the check
- *-move* option moves corrupted files to */lost+found*
 - A corrupted file is one where all replicas of a block are missing
- *-delete* option deletes corrupted files

dfsadmin

- The *hdfs dfsadmin* command provides a number of administrative features including:
- List information about HDFS on a per-datanode basis:
 - *hdfs dfsadmin -report*
- Re-read the `dfs.hosts` and `dfs.hosts.exclude` files
 - These are defined in `hdfs-site.xml`, contains file of list of hosts which are (not) allowed to connect to namenode
 - *hdfs dfsadmin -refreshNodes*

Cluster Rebalancing

- An HDFS cluster can become ‘unbalanced’
 - Some nodes have much more data on them than others
 - Example: add a new node to the cluster
 - Even after adding some files to HDFS, this node will have far less data than the others
 - During MapReduce processing, this node will use much more network bandwidth as it retrieves data from other nodes
- Clusters can be rebalanced using the *hdfs balancer* utility

Using hdfs balancer

- *hdfs balancer* reviews data block placement on nodes and adjusts blocks to ensure all nodes are within x% utilization of each other
 - Utilization is defined as amount of data storage used
 - x is known as the threshold
- A node is under-utilized if its utilization is less than (average utilization - threshold)
- A node is over-utilized if its utilization is more than (average utilization + threshold)
- Note: hdfs balancer does not consider block placement on individual disks on a node
 - Only the utilization of the node as a whole

Using hdfs balancer

- Syntax:

hdfs balancer -threshold x

- Threshold is optional

- Defaults to 10 (i.e., 10% difference in utilization between nodes)

- Rebalancing can be canceled at any time

- Interrupt the command with Ctrl+C

When to Rebalance

- Rebalance immediately after adding new nodes to the cluster
- Rebalance during non-peak usage times
 - Rebalancing does not interfere with any existing MapReduce jobs
 - However, it does use bandwidth

Job Management

- To view all jobs running on the cluster
 - *mapred job -list*
- To view all jobs including completed ones
 - *mapred job -list all*

```
15/08/22 17:03:26 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
```

```
Total jobs:7
```

Priority	JobId	State	StartTime	UserName	Queue
dMem	UsedContainers	RsvdContainers	UsedMem	RsvdMem	Neede
NORMAL	AM info				
N/A	job_1434693510565_0003	SUCCEEDED	1439364547526	hduser	default
N/A	http://localhost:8088/proxy/application_1434693510565_0003/	N/A	N/A	N/A	
N/A	job_1434693510565_0004	SUCCEEDED	1439447100049	hduser	default
N/A	http://localhost:8088/proxy/application_1434693510565_0004/	N/A	N/A	N/A	
N/A	job_1434693510565_0005	SUCCEEDED	1440045711153	hduser	default
N/A	http://localhost:8088/proxy/application_1434693510565_0005/	N/A	N/A	N/A	
N/A	job_1434693510565_0006	SUCCEEDED	1440047295680	hduser	default
N/A	http://localhost:8088/proxy/application_1434693510565_0006/	N/A	N/A	N/A	

Display individual job status

- To display individual job status:
 - *mapred job -status <job_id>*
 - It provides completion percentage, values of counters, etc
 - Job name is not displayed
 - The Web user interface is the most convenient way to view more details about an individual job

Kill a job

- It is important to note that once a user has submitted a job, they can not stop it just by hitting CTRL+C on their terminal
 - This stops job output appearing on the user's console
 - The job is still running on the cluster!

Kill a job

- To kill a job use *mapred job -kill <job_id>*

```
[training@localhost ~]$ mapred job -list
1 jobs currently running
JobId    State    StartTime      UserName       Priority       SchedulingInfo
job_201110311158_0009  1        1320210791739  training      NORMAL      NA

[training@localhost ~]$ mapred job -kill job_201110311158_0009
Killed job job_201110311158_0009

[training@localhost ~]$ mapred job -list
0 jobs currently running
JobId    State    StartTime      UserName       Priority       SchedulingInfo
```

Web monitoring

- Namenode Web UI:
 <NameNode ADDR>:50070
- JobTracker Web UI (Hadoop 1.x):
 <JobTracker ADDR>:50030
- ResourceManager Web UI (YARN):
 <ResourceManager ADDR>:8088

Hadoop Configuration Files

- Each machine in the Hadoop cluster has its own set of configuration files
- Configuration files all reside in Hadoop's conf directory
 - Typically `/etc/hadoop/conf`
- Most of the configuration files are written in XML
- Upon startup, the Hadoop daemons access the configuration files
 - After modifying configuration parameters, you must restart Hadoop daemons for your changes to take effect

Hadoop Configuration Files Overview

File	Type of Configuration
<code>core-site.xml</code>	Core
<code>hdfs-site.xml</code>	HDFS
<code>mapred-site.xml</code>	MapReduce
<code>hadoop-policy.xml</code>	Access control policies
<code>log4j.properties</code>	Logging
<code>hadoop-metrics.properties</code> , <code>hadoop-metrics2.properties</code>	Metrics
<code>include</code> , <code>exclude</code> (file names are configurable)	Host inclusion/exclusion in a cluster
<code>allocations.xml</code> (file name is configurable)	FairScheduler
<code>masters</code> , <code>slaves</code>	Scripted startup (not recommended)
<code>hadoop-env.sh</code>	Environment variables

Configuration Value Precedence

- Configuration parameters can be specified more than once
- Highest precedence value takes priority
- Precedence order (lowest to highest):
 - *-site.xml on the slave node
 - *-site.xml on the client machine
 - Values set explicitly in the Job object for a MapReduce job

Configuration Value Precedence

- If a value in a configuration file is marked as final it overrides all others

```
<property>  
  <name>some.property.name</name>  
  <value>somevalue</value>  
  <final>true</final>  
</property>
```

Important Configurations

- core-site.xml:

<code>hadoop.tmp.dir</code>	<p>Base temporary directory, both on the local disk and in HDFS. Default is <code>/tmp/hadoop-\${user.name}</code>. Used by all nodes.</p> <p>[This parameter is used to derive defaults for numerous other configuration parameters. For example, the default value for <code>dfs.data.dir</code> is <code>file://\${hadoop.tmp.dir}/dfs/name</code></p>
-----------------------------	---

- In our system, it's `/tmp/hadoop-hduser/`

core-site.xml

`fs.default.name`

The name of the default filesystem. Usually includes the file system type, plus the NameNode's hostname and port number. Example:
`hdfs://<your_namenode>:8020/`
Used by every machine which needs access to the cluster, including all nodes running Hadoop daemons.

- In YARN, it's replaced by `fs.defaultFS`

hdfs-site.xml

- The single most important configuration value on your entire cluster, used by the NameNode
- **dfs.name.dir**: A comma separated list of directories, describing where namenode stores the HDFS metadata (fsimage + Edit log)
 - default value = `${hadoop.tmp.dir}/dfs/name`
- Loss of a NameNode's metadata will result in the loss of all the data in its namespace
 - Although the blocks will remain, there is no way of reconstructing the original files without the metadata
- There must be at least two disks (or a RAID volume) on the NameNode, plus an NFS mount elsewhere on the network
 - Failure to set this correctly will result in eventual loss of your cluster's data

hdfs-site.xml

- A NameNode will write to the edit log in all directories in *dfs.name.dir* synchronously
- If a directory in the list disappears, the NameNode will continue to function
 - It will ignore that directory until it is restarted
- Note: no space between the comma and next directory name in the list!
 - Example: */disk1/dfs/nn,/disk2/dfs/nn*

hdfs-site.xml

- ❑ **dfs.block.size**: The block size for new files, in bytes.
 - ❑ Default is 67108864 (64MB)
- ❑ **dfs.data.dir**: A comma separated list of directories, describing where a datanode stores its blocks
 - ❑ default value = `${hadoop.tmp.dir}/dfs/data`
 - ❑ No space between the comma and the path
 - ❑ Round-robin writes to the directories in the list
 - ❑ Used by DataNodes
 - ❑ Can be different on each DataNode

hdfs-site.xml

- ❑ **fs.checkpoint.dir:** property for secondary namenode to store its checkpoints for the filesystem
 - ❑ default value = `${hadoop.tmp.dir}/dfs/namespacesecondary`
- ❑ You can see `fsimage`, `edits` files and the md5 CRC code in *dfs.data.dir* and *fs.checkpoint.dir*

hdfs-site.xml

- ❑ **dfs.http.address:** The address and port used for the NameNode Web UI, used by NameNode
 - ❑ Default is `<your_namenode>:50070`
- **dfs.replication:** The number of times each block should be replicated when a file is written.
 - ❑ Default is 3

Environment Setup: `hadoop-env.sh`

- `hadoop-env.sh` sets environment variables necessary for Hadoop to run
 - `HADOOP_CLASSPATH`
 - `HADOOP_HEAPSIZE`
 - `HADOOP_LOG_DIR`
 - `HADOOP_PID_DIR`
 - `JAVA_HOME`
- Values are sourced into all Hadoop control scripts and therefore the Hadoop daemons
- If you need to set environment variables, do it here to ensure that they are passed through to the control scripts